B.TECH ARTIFICIAL INTELLIGENCE & MACHINE LEARNING · CA2 CAPSTONE PROJECT

Decision Debt OS

Turning business decisions into measurable learning



Name: [Student Name] **PRN:** [PRN Number] **Branch:** [Branch]

Division: [Division] **Institution:** [Institution Name] **Academic Year:** 2026–27

THE PROBLEM

Decisions get made. Assumptions get forgotten.



Dashboards and BI tools show what happened. Project trackers show what is scheduled. Neither one records **what a team believed would happen** before they acted — so when the belief turns out to be wrong, nothing captures that it was ever tested, and the same unexamined assumption drives the next decision.

CORE INSIGHT

A decision is not finished when it is made.

> *“A business decision is incomplete until its original assumption is compared with the real-world outcome.”*

This gap has a name:

Decision Debt — the accumulation of decisions whose assumptions were never checked against reality.

DEFINING THE PROBLEM

What Is Decision Debt?

Unresolved Assumptions

A decision was made, but its underlying belief was never written down in a checkable form.

Overdue Reviews

A deadline for checking the outcome has passed, and no one has looked at the result.

Invalidated Assumptions Still Driving Work

The assumption was checked, found false, but the decision or policy it justified is still active.

Repeated Forecasting Errors

The same type of assumption keeps failing across different decisions, and no one notices the pattern.

THE SOLUTION

Capture → Structure → Forecast → Measure → Evaluate → Learn



UNSTRUCTURED INPUT

"Let's offer the 10% discount to high-risk customers — I think it'll cut churn to around 18% within a month."

STRUCTURED DECISION RECORD

INTERVENTION	10% retention discount
TARGET METRIC	30-day churn, high-risk cohort
FORECAST	25% → 18%
CONFIDENCE	62%
DEADLINE	+30 days

WORKED EXAMPLE

One Decision, Fully Specified

> *"We believe a 10% retention discount will reduce 30-day churn among high-risk customers from 25% to 18%."*

OWNER

Retention Lead, Growth Team

CONFIDENCE (FORECAST)

62% self-reported probability

INTERVENTION

10% discount, auto-applied at checkout

TARGET METRIC

30-day churn rate, high-risk cohort

DEADLINE

30 days from decision date

ESTIMATED COST

~$4.80 / customer (10% of avg. order)

KILL CONDITION

If interim churn > 25% at day 15, stop offer

ACTUAL RESULT

Pending — 30-day window not yet elapsed

Baseline shown here (25%) is one risk-tier definition; the dataset's observed high-risk cohort baseline is 37% — see Slide 7.

DATASET & PIPELINE — SYNTHETIC, FOR DEMONSTRATION ONLY

From Customer Record to Decision Evaluation

Customer profile → Churn-risk prediction → Retention decision → Treatment / control group → 30-day outcome → Decision evaluation

SYNTHETIC CUSTOMER-RETENTION DATASET

RECORDS 10,000 fictional customer rows

FIELDS age, gender, country, income, spending score, purchases, avg. order value, membership years, site visits/mo, cart-abandon rate, churn flag, feedback text, last purchase date

STATUS Clearly labelled synthetic data — used for demonstration only; not a real company's data

Real, observed churn rate by risk segment

